

Hasini Gangoni

Data Scientist

Jersey City, NJ • 862-822-9670 • hasinigangoni01@gmail.com

PROFESSIONAL SUMMARY

Data Scientist with 2+ years of experience applying statistical modeling, machine learning, and advanced analytics to drive data-driven decision-making. Proficient in Python (pandas, NumPy, scikit-learn) and SQL for end-to-end analytics — from data preparation and feature engineering to model validation and business interpretation. Skilled at translating complex analytical findings into clear, actionable insights for cross-functional stakeholders in enterprise environments.

EXPERIENCE

Data Scientist | Merck

June 2025 – Present

- Developed and applied statistical and machine learning models (regression, classification) to support business forecasting and decision-making for enterprise clients.
- Executed end-to-end analytics lifecycle including data preparation, exploratory analysis, feature engineering, model validation, and result interpretation.
- Partnered with business stakeholders to define analytical questions and success metrics, translating model outputs into clear, actionable recommendations.
- Documented model methodologies and assumptions, contributing to shared best practices and reproducible analytical workflows.

Data Engineer / Applied Analyst | Accenture

Aug 2022 – May 2025

- Prepared and transformed large-scale structured datasets using Python (pandas, NumPy) and SQL to support ML model training and statistical analysis pipelines.
- Optimized complex SQL queries for data extraction and aggregation across enterprise-scale relational databases, improving processing efficiency by 30%.
- Built and maintained data pipelines on AWS (S3, Redshift, Glue) to deliver clean, reliable datasets for analytics and data science workflows.
- Collaborated with data scientists to design feature engineering pipelines and validate model inputs using Apache Spark and Hive.
- Automated data quality checks and lineage tracking using Airflow and DBT, ensuring accuracy and trustworthiness of analytical outputs.
- Integrated data from diverse sources (APIs, RDBMS, Kafka) into unified data models, enabling consistent cross-functional reporting and analysis.

EDUCATION

M.S. in Information Systems — Pace University, Brooklyn, NY

Sep 2022 – May 2024

B.Tech in Electrical & Electronics Engineering — Vishnu Raju Institute of Technology, Hyderabad

Aug 2018 – Jul 2022

TECHNICAL SKILLS

Machine Learning: Regression, classification, time-series forecasting, anomaly detection — scikit-learn, model validation & interpretation

Programming: Python (pandas, NumPy, scikit-learn), SQL, Java, Scala — end-to-end analytics and automation

Statistical Analysis: Exploratory data analysis, feature engineering, A/B testing, experiment design and result interpretation

Data Platforms: AWS Redshift, S3, Glue, EMR — enterprise-scale structured and semi-structured data processing

Big Data & Pipelines: Apache Spark, Hadoop, Hive, Airflow, DBT — batch and streaming pipeline design

Databases: PostgreSQL, MySQL, MongoDB, Snowflake — relational and NoSQL data modeling

Cloud Platforms: AWS (S3, Redshift, EC2, Glue, EMR, RDS), GCP, Azure

Tools: Git, Jenkins, Linux/Unix, Kafka, Looker, Tableau

CERTIFICATIONS

- AWS Certified Data Analytics – Specialty
- Databricks Certified Spark Developer
- Oracle Certified Java Programmer
- Python for Data Analysis (pandas, NumPy) — Udemy