

Hasini Gangoni

Data Engineer

Jersey City, NJ • 862-822-9670 • hasinigangoni01@gmail.com

PROFESSIONAL SUMMARY

Data Engineer with 2+ years of experience designing and maintaining scalable data pipelines, ETL workflows, and AI/ML platform infrastructure. Proficient in Python, Apache Spark, Kafka, and Airflow with hands-on experience processing high-volume streaming data and building CI/CD pipelines. Skilled in cloud platforms (AWS, GCP, Azure), NoSQL and relational databases, and Infrastructure as Code tooling. Adept at collaborating with data scientists to operationalize ML features and delivering low-latency, high-throughput data solutions in enterprise environments.

EXPERIENCE

Data Engineer | Merck

June 2025 – Present

- Develop, deploy, and maintain scalable ETL and streaming data pipelines using Apache Spark Streaming and Kafka to ingest and process high-velocity data from enterprise sources, enabling near-real-time analytics.
- Build and maintain CI/CD pipelines using GitHub Actions to ensure code quality, automated testing, and rapid deployment of data software across environments.
- Implement and manage AI/ML feature engineering pipelines, collaborating with data scientists to operationalize features for model training and inference on AWS (Amazon Bedrock, SageMaker).
- Automate infrastructure provisioning and orchestration for data and AI workloads using Terraform and CloudFormation (Infrastructure as Code).
- Monitor, troubleshoot, and optimize data pipelines to ensure low-latency, high-throughput data flows across cloud and on-premise environments.
- Design and maintain data workflows in Apache Airflow for orchestrating complex multi-step ETL jobs across distributed systems.

Data Engineer | Accenture

Aug 2022 – May 2025

- Built scalable batch and streaming data pipelines using Apache Spark, Kafka, and Java to process 2TB+ of data daily across distributed cloud infrastructure.
- Designed and maintained ETL workflows integrating data from APIs, RDBMS (PostgreSQL, MySQL), NoSQL (MongoDB), and Kafka topics into unified data models.
- Optimized complex SQL queries and data retrieval strategies in PostgreSQL and AWS Redshift, improving processing efficiency by 30%.
- Implemented data quality checks and lineage tracking using Airflow and DBT, ensuring pipeline reliability and reproducibility.
- Managed source control and version workflows using GitHub; contributed to team CI/CD practices for data pipeline deployments.
- Built and maintained data lakes on AWS (S3, Redshift, Glue, EMR) and participated in migration of on-premise ETL jobs to cloud infrastructure.
- Collaborated with data scientists to prepare clean, feature-engineered datasets for ML model training on Linux-based environments.

EDUCATION

M.S. in Information Systems — Pace University, Brooklyn, NY

Sep 2022 – May 2024

B.Tech in Electrical & Electronics Engineering — Vishnu Raju Institute of Technology, Hyderabad

Aug 2018 – Jul 2022

TECHNICAL SKILLS

Streaming & Messaging: Apache Kafka, Spark Streaming — real-time ingestion and processing of high-velocity data

Data Pipelines & ETL: Apache Spark, Hadoop, Hive, Airflow, DBT — batch and streaming ETL design and orchestration

Programming: Python, Java, Scala, SQL — pipeline development, automation, and scripting on Linux/Unix systems

Databases: PostgreSQL, MySQL, MongoDB (NoSQL), Snowflake, AWS Redshift — relational and NoSQL data modeling

AI/ML Platforms: Amazon Bedrock, AWS SageMaker — ML model management, feature pipelines, and inference support

CI/CD & IaC: GitHub Actions, Terraform, CloudFormation — pipeline automation and infrastructure provisioning

Cloud Platforms: AWS (S3, Redshift, Glue, EMR, EC2, RDS, SageMaker, Bedrock), GCP, Azure

Tools & OS: Git, GitHub, Jenkins, Linux/Unix, Docker

CERTIFICATIONS

- AWS Certified Data Analytics – Specialty
- Databricks Certified Spark Developer
- Oracle Certified Java Programmer
- Python for Data Analysis (pandas, NumPy) — Udemy